

Assignment 3

Statistics and Data Science 365/565

Due: October 9 (before midnight)

I worked with Jackson Simon on Problem 1, Part 2 and Problem 2, Part B.

This homework treats variable selection and shrinkage. Problem 1 is a conceptual question, and Problems 2 and 3 are applied problems. Throughout this assignment, remember that for a vector $\beta \in \mathbb{R}^p$,

$$\|\beta\|_2 = \sqrt{\sum_{j=1}^p \beta_j^2}$$
$$\|\beta\|_1 = \sum_{j=1}^p |\beta_j|$$

Problem 1 (25 points) LASSO Risk

Suppose we estimate β in a linear regression model by minimizing

$$\widehat{\beta}^{(s)} = \arg \min_{\beta \in \mathbb{R}^{p+1}} \|Y - X\beta\|_2^2 \quad \text{subject to} \quad \|\beta\|_1 \leq s$$

for a particular value of s .

Part 1.

For (a) through (e), indicate which of (i) through (v) is correct, and justify your answer.

Part a.

As we increase s from 0, the training RSS of $\widehat{\beta}^{(s)}$ will:

- (i) increase initially, and then eventually start decreasing in an inverted U shape
- (ii) decrease initially, and then eventually start increasing in a U shape
- (iii) steadily increase
- (iv) steadily decrease
- (v) remain constant

Solution: (iv) steadily decrease As s increases from zero, we are allowing the size of our β_j coefficients to increase. Starting at s zero, we are only using the intercept as our $\widehat{\beta}$ and forcing other coefficients to be zero. Adding more parameters to our model by allowing the other coefficients to be nonzero, we are reducing RSS. With training data, adding more and more parameters into our model will continuously decrease RSS

Part b.

Repeat (a) for the test RSS of $\hat{\beta}^{(s)}$.

Solution: (ii) decrease initially, and then eventually start increase in a U shape As stated above, s increasing from zero, we are adding more parameters to our model by allowing the other coefficients to be nonzero, we are reducing RSS. However, since we fit our model to the training data, increasing s and allowing too many parameters we will begin to overfit our test data and RSS will begin to increase again. Think of how 1-NN fits the training data perfectly but overfits the test data.

Part c.

Repeat (a) for the variance of $\hat{\beta}^{(s)}$.

Solution: (iii) steadily increase As our model becomes more complex (i.e. more predictor variables used), variance increases. So as s increases and we include more parameters in our model, variance will steadily increase.

Part d.

Repeat (a) for the squared bias of $\hat{\beta}^{(s)}$.

Solution: (iv) steadily decrease Contrary to above, as our model becomes more complex, squared bias decreases. So as s increases and we include more parameters in our model, squared bias will steadily decrease. For s approaching infinity (or $\lambda=0$), then we have an unbiased estimator for the OLS model.

Part e.

Repeat (a) for the irreducible error of $\hat{\beta}^{(s)}$.

Solution: (v) remain constant As the name suggests, the irreducible error is thus irreducible error. Our data is inherently noisy as thus is modeled as $y_i = f(X_i) + \epsilon_i$ where ϵ_i is irreducible error. No matter what function we choose for $f(X_i)$ (i.e. what parameters and their coefficients we choose to include by changing the values of s), these errors will remain constant.

Part 2.

Now we will compute the lasso estimates on simulated data. That is, we will find the minimizer

$$\hat{\beta}^{(\lambda)} = \arg \min_{\beta \in \mathbb{R}^{p+1}} \|Y - X\beta\|_2^2 + \lambda \|\beta\|_1$$

where λ is some fixed constant.

Part a.

What do you think is the relationship between s in Part 1 and λ ? As we increase λ from 0, what will happen to the bias and variance of our estimator $\hat{\beta}^{(\lambda)}$?

Solution: Comparing this equation to the one in Part 1, we can see that large values of λ correspond to small values of s . As we increase the weight of the penalty (λ), we want to decrease the size of the beta

coefficients (s) to compensate for the penalty increase and continue to minimize RSS. Thus as we increase λ from 0, we will be including less and less parameters in our model according to the following:

$$\hat{\beta} = \begin{cases} Y - \lambda, & \text{if } Y > \lambda \\ 0, & \text{if } Y \leq |\lambda| \\ Y + \lambda, & \text{if } Y < -\lambda \end{cases}$$

Therefore model complexity will decrease, thus variance will decrease and squared bias will increase (the opposite of what we see when s increases).

Part b.

Generate independent predictors, X_1, X_2, X_3, X_4 , and store them in the $n \times 4$ matrix X , as follows:

```
n <- 1000
p <- 4
X <- matrix(rnorm(n*p, mean=1), nrow = n, ncol = p)
```

Also generate an outcome vector $Y \in \mathbb{R}^n$ according to the model:

$$Y_i = 17 + (.005) * X_{i1} + 117 * X_{i2} + .6 * X_{i3} + 52 * X_{i4} + \epsilon_i$$

where the $\epsilon_i \sim N(0, 1)$.

For various values of λ (between 0 and 0.1), fit a lasso model (using the `glmnet` package) to the generated data, and store each fitted lasso estimate $\hat{\beta}^{(\lambda)}$. For B times, repeat the process of generating data and, for the same set of values for lambda, fitting a lasso and storing the result. (It is up to you to decide how big B should be. i.e. How many repetitions to do.) To estimate the bias and variance, let

$$\mathbb{E}\hat{\beta}^{(\lambda)} = \frac{1}{B} \sum_{b=1}^B \hat{\beta}^{(b,\lambda)}$$

where $\hat{\beta}^{(b,\lambda)}$ represents the vector of lasso coefficients on round b and penalization parameter λ .

Estimate the squared bias via:

$$\widehat{\text{Bias}}^2(\hat{\beta}^{(\lambda)}) = \|\mathbb{E}\hat{\beta}^{(\lambda)} - \beta\|_2^2$$

And estimate the variance via:

$$\widehat{\text{Var}}(\hat{\beta}^{(\lambda)}) = \frac{1}{B} \sum_{b=1}^B \|\hat{\beta}^{(b,\lambda)} - \mathbb{E}\hat{\beta}^{(\lambda)}\|_2^2$$

Plot the estimated bias and variance for each λ . For a large enough number of repetitions B , you should see a clear pattern. Summarize your results and compare them to Part (a).

#Generate our y values, we are going to need to do this B times, lets just set B to 100 first.

```
B <- 1000
lambdas <- seq(0,0.1,.005)
coefficients <- array(data = rep(0,B*length(lambdas)*4), dim =c(B,length(lambdas),4))
beta_vec <- as.vector(c(.005,117,.6,52))
```

```
for (i in 1:B) {
  errors <- rnorm(1000)
  y <- 17 + X%*%beta_vec + errors
  mod <- glmnet(X, y, family="gaussian", standardize=TRUE, alpha=1, lambda=lambdas)
  for (j in c(1:length(lambdas))) {
    coefficients[i,j,] <- coef(mod)[2:5,j]
  }
}
```

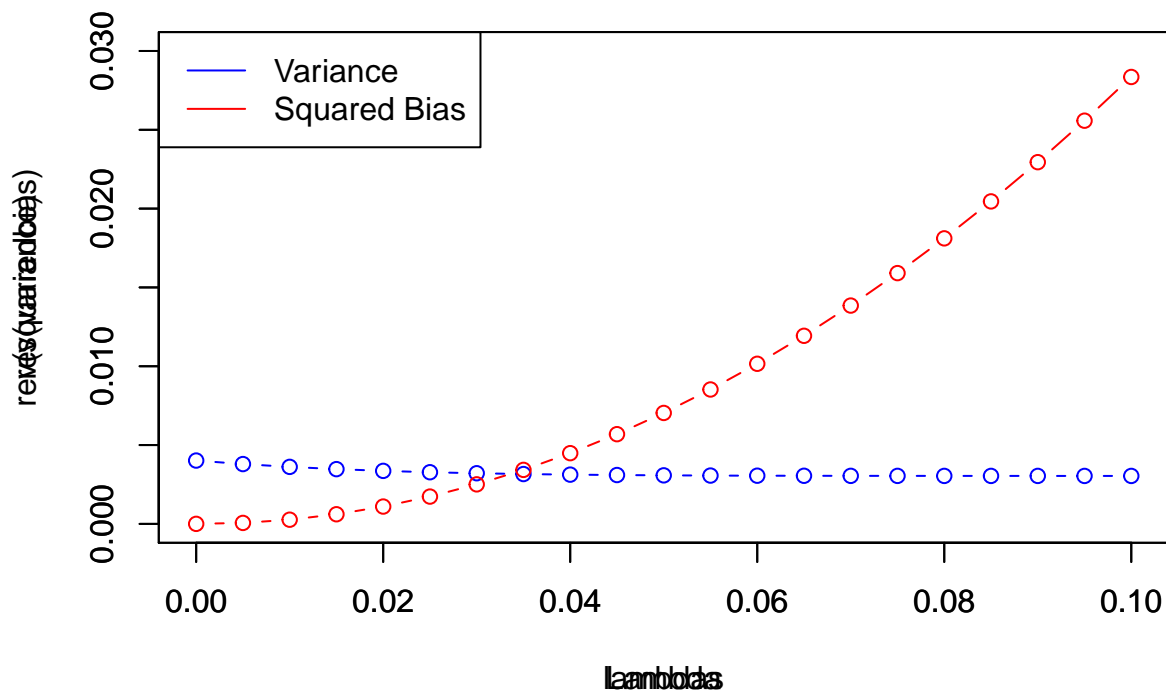
```

squaredbias <- rep(0,length(lambdas))
variance <- rep(0,length(lambdas))
expectations <- array(data = rep(0,length(lambdas)*4), dim = c(length(lambdas),4))
for (i in 1:length(lambdas)) {
  expectations[i,] <- colMeans(coefficients[,i,])
  squaredbias[i] <- sum((expectations[i,] - beta_vec)^2)
  variance[i] <- sum(colMeans(t(apply(coefficients[,i,], 1, function(x) (x-expectations[i,])^2))))
}

plot(lambdas,rev(variance),col="blue",type="b",ylim=c(0,.03),main="Squared Bias and Variance by Lambda")
par(new=T)
plot(lambdas,rev(squaredbias),ylim=c(0,.03),col="red",type="b",xlab="Lambda")
legend("topleft",legend=c("Variance", "Squared Bias"),col=c("blue","red"),lty=1)

```

Squared Bias and Variance by Lambda



Solution: According to my results here, squared bias increase and variance decreases as lambda increases (s decreases). This is in agreement with what I argue in part 1, that variance constantly increases and squared bias constantly decreases.

Problem 2 (15 points) Boston Housing Values

The value of homes in Boston is a dataset included in the `MASS` library in R with the name `Boston`. You may use the `glmnet` package to fit the lasso and ridge regression models.

Part a.

Start by doing some data exploration. What are n and p in this dataset? Are there any missing data points, outliers, or explanatory variables that seem highly correlated with each other? What variables seem most

closely associated with the outcome variable, medv? You are not limited to these questions; examine whatever you think is interesting and reasonable. Show plots, accompanied with comments, that illustrate your findings.

```
dim(Boston)

## [1] 506 14

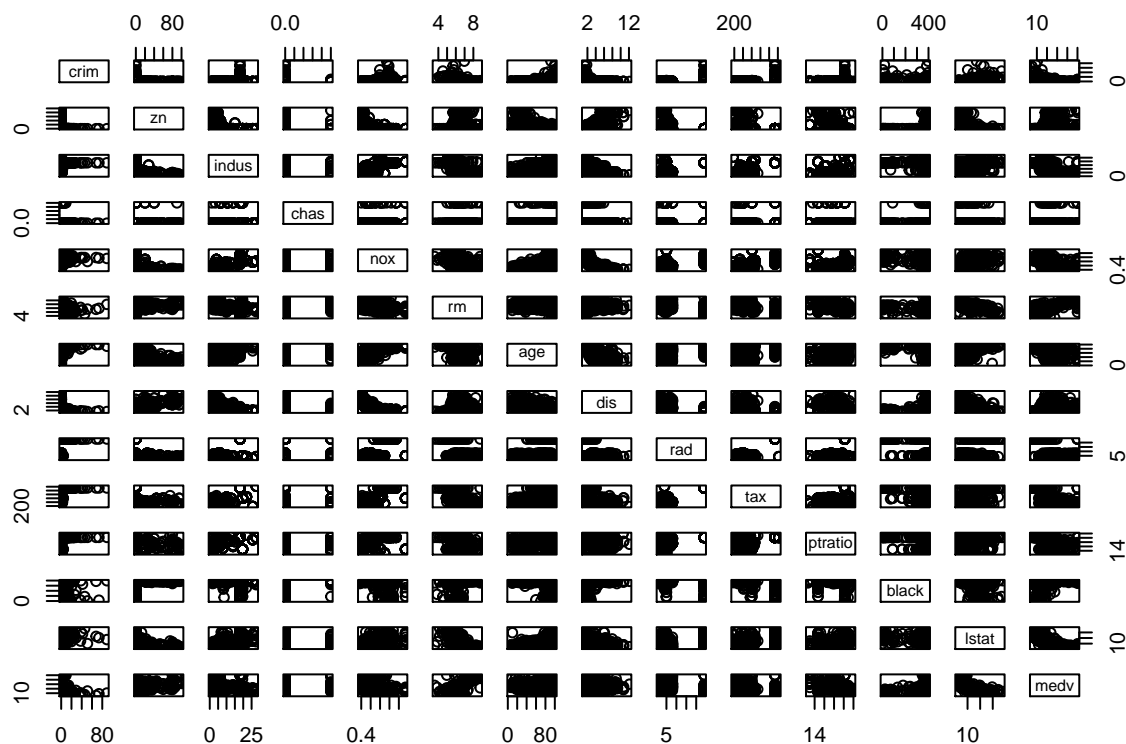
summary(is.na(Boston))

##      crim      zn      indus      chas
## Mode :logical Mode :logical Mode :logical Mode :logical
## FALSE:506    FALSE:506    FALSE:506    FALSE:506
##      nox      rm      age      dis
## Mode :logical Mode :logical Mode :logical Mode :logical
## FALSE:506    FALSE:506    FALSE:506    FALSE:506
##      rad      tax      ptratio    black
## Mode :logical Mode :logical Mode :logical Mode :logical
## FALSE:506    FALSE:506    FALSE:506    FALSE:506
##      lstat      medv
## Mode :logical Mode :logical
## FALSE:506    FALSE:506

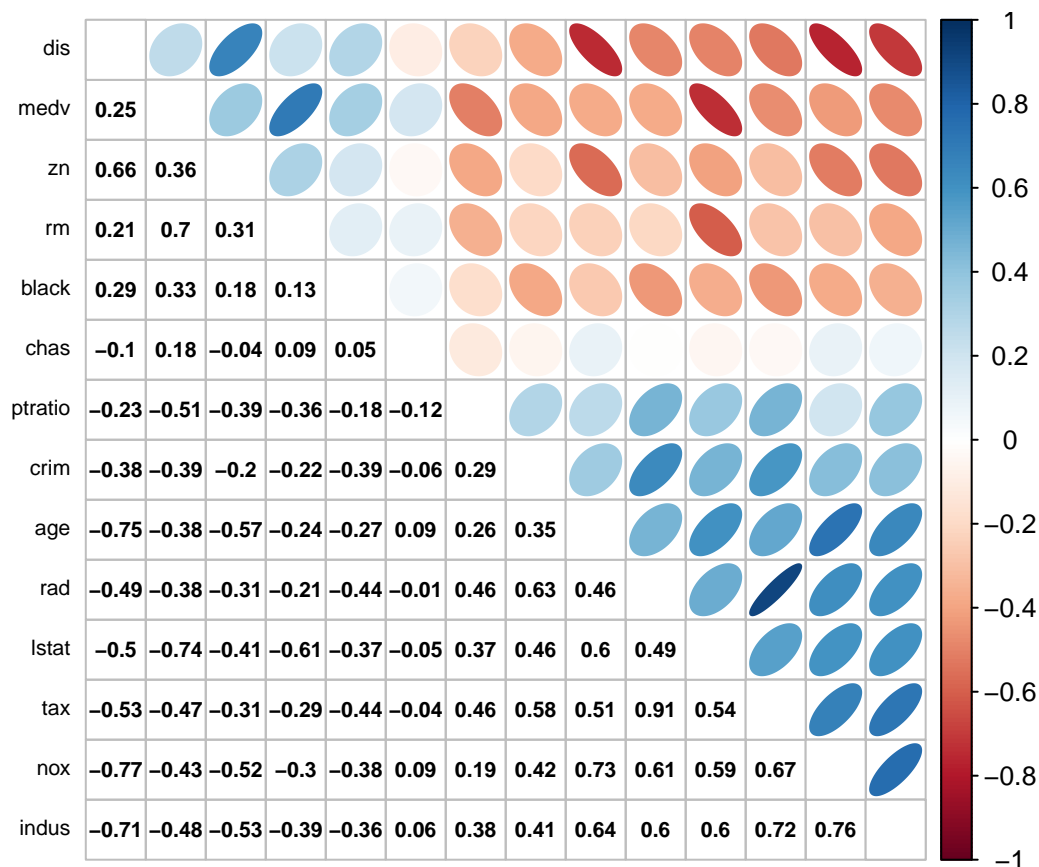
summary(Boston)

##      crim      zn      indus      chas
## Min.   : 0.00632 Min.   : 0.00 Min.   : 0.46 Min.   :0.00000
## 1st Qu.: 0.08204 1st Qu.: 0.00 1st Qu.: 5.19 1st Qu.:0.00000
## Median : 0.25651 Median : 0.00 Median : 9.69 Median :0.00000
## Mean   : 3.61352 Mean   : 11.36 Mean   :11.14 Mean   :0.06917
## 3rd Qu.: 3.67708 3rd Qu.: 12.50 3rd Qu.:18.10 3rd Qu.:0.00000
## Max.   :88.97620 Max.   :100.00 Max.   :27.74 Max.   :1.00000
##      nox      rm      age      dis
## Min.   :0.3850 Min.   :3.561 Min.   : 2.90 Min.   : 1.130
## 1st Qu.:0.4490 1st Qu.:5.886 1st Qu.:45.02 1st Qu.: 2.100
## Median :0.5380 Median :6.208 Median :77.50 Median : 3.207
## Mean   :0.5547 Mean   :6.285 Mean   :68.57 Mean   : 3.795
## 3rd Qu.:0.6240 3rd Qu.:6.623 3rd Qu.:94.08 3rd Qu.: 5.188
## Max.   :0.8710 Max.   :8.780 Max.   :100.00 Max.   :12.127
##      rad      tax      ptratio    black
## Min.   : 1.000 Min.   :187.0 Min.   :12.60 Min.   : 0.32
## 1st Qu.: 4.000 1st Qu.:279.0 1st Qu.:17.40 1st Qu.:375.38
## Median : 5.000 Median :330.0 Median :19.05 Median :391.44
## Mean   : 9.549 Mean   :408.2 Mean   :18.46 Mean   :356.67
## 3rd Qu.:24.000 3rd Qu.:666.0 3rd Qu.:20.20 3rd Qu.:396.23
## Max.   :24.000 Max.   :711.0 Max.   :22.00 Max.   :396.90
##      lstat      medv
## Min.   : 1.73 Min.   : 5.00
## 1st Qu.: 6.95 1st Qu.:17.02
## Median :11.36 Median :21.20
## Mean   :12.65 Mean   :22.53
## 3rd Qu.:16.95 3rd Qu.:25.00
## Max.   :37.97 Max.   :50.00

plot(Boston)
```



```
cor1 <- round(cor(Boston, use = "pairwise.complete.obs"),2)
corrplot.mixed(cor1,lower.col="black", upper = "ellipse", tl.col = "black", number.cex=.7,
               order = "FPC", tl.pos = "lt", tl.cex=.7, sig.level = .05)
```



Comments: By using the `dim()` function, we see that `n` is 506 (number of observations) and `p` is 14 (number of predictor variables). Additionally, there are no missing values in the dataset. Then using the `corrplot.mixed` function from the `corrplot` library, we see that our strongest associations between outcome variable `medv` and the predictor variables is a 0.7 R-squared with the numbers of rooms (`rm`) and a -0.74 correlation with the percent in the “lower status” of the populations (`lstat`). These seem to make sense, I am surprised to see that there is a strong negative correlation (-0.51) with the pupil-to-teacher ratio in local schools (`ptratio`). It is worth noting that many predictor variables are correlated with each other. For example, `tax` and `rad` (index of highway accessibility) are strongly positively correlated (0.91) while `dis` and `indus` are strongly negatively correlated (-0.71), this second example suggests that the largest employment centers are retail oriented. Observing the summary data for the Boston housing data, I do not see many extreme values.

Part b.

Randomly sample 100 rows of the Boston dataset to use as the test set, and use the rest for training. Fit lasso and ridge regression models to predict `medv` with all other variables as predictors. For various values of λ , use cross-validation to select an optimal values to use in the lasso and ridge regression models. (You may use the `cv.glmnet()` in the `glmnet` package to do the cross-validation.) Also, plot the model coefficients against the values of λ . How does each coefficient change as λ increases in the two models? How is this behavior similar or different between the lasso and ridge regression approaches?

```
set.seed(1)
testrows <- sample(nrow(Boston),100)
test <- Boston[testrows,]
train <- Boston[-testrows,]

#lasso
cvfit <- cv.glmnet(x=as.matrix(train[-14]),y=as.matrix(train[14]),standardize=TRUE,nfold=10,alpha=1)
cvfit$lambda.min

## [1] 0.01773275

predict(cvfit, newx = as.matrix(test[-14]), s = "lambda.min")

##          1
## 135 13.428197
## 188 33.042812
## 289 27.138334
## 457 12.834164
## 102 25.373309
## 451 16.752445
## 473 21.595768
## 330 24.817806
## 314 25.256494
## 31  11.916388
## 103 20.328959
## 88  25.636672
## 340 21.055558
## 190 34.347536
## 379 15.280124
## 245 16.865574
## 352 21.780879
## 486 21.466260
## 186 24.534903
## 492 14.185122
```

455 15.619387
496 16.729627
316 20.426500
61 18.003028
129 18.720548
488 20.333664
7 23.244214
184 30.484882
416 10.114886
163 40.626795
230 30.587811
285 31.672446
234 36.974534
89 30.635964
391 16.238724
315 25.381856
374 5.456154
51 21.599244
339 21.934629
193 33.138876
383 12.588280
301 31.036182
364 19.730106
257 36.992520
491 4.478407
464 21.573125
11 19.827419
220 29.788818
336 20.652702
317 17.811675
218 28.039414
392 16.498874
199 35.018057
111 20.466985
32 18.216110
45 23.120256
143 15.112473
233 37.862886
297 27.568242
182 27.203705
408 18.843133
131 19.806944
204 41.163340
148 8.619078
288 26.912300
114 20.730302
211 22.849600
337 20.052932
37 22.018248
466 17.087189
443 17.876000
366 12.547373
151 20.706581
145 8.787821


```

## 206 22.383808
## 385 2.826338
## 372 23.659503
## 168 23.147581
## 333 23.712177
## 411 15.328508
## 481 22.552406
## 303 28.644318
## 170 26.538121
## 138 19.135789
## 320 21.323181
## 86 27.740210
## 299 29.258463
## 469 16.382184
## 485 18.655429
## 60 21.011262
## 100 32.287963
## 25 15.919708
## 266 27.458186
## 362 18.151930
## 321 24.743018
## 328 19.534411
## 187 35.610790
## 429 14.171556
## 331 22.382446
## 247 20.464967

#ridge
cvfit2 <- cv.glmnet(x=as.matrix(train[-14]),y=as.matrix(train[14]),standardize=TRUE,nfold=10,alpha=0)
cvfit2$lambda.min

## [1] 0.7499609

predict(cvfit2, newx = as.matrix(test[-14]), s = "lambda.min")

## 1
## 135 14.065919
## 188 32.439413
## 289 26.809812
## 457 12.943005
## 102 25.473553
## 451 16.811247
## 473 21.074537
## 330 25.551243
## 314 25.012835
## 31 12.884089
## 103 20.385685
## 88 25.219023
## 340 21.154120
## 190 33.963301
## 379 15.385854
## 245 17.400515
## 352 23.341650
## 486 21.133595
## 186 24.163133
## 492 15.580229

```

455 15.817688
496 17.378313
316 20.512539
61 18.267057
129 19.211275
488 19.756576
7 23.444163
184 29.565731
416 10.536844
163 40.018625
230 29.810279
285 31.685748
234 36.373414
89 30.373916
391 16.031372
315 25.337309
374 5.734321
51 21.822012
339 22.003145
193 33.214111
383 12.613615
301 31.015032
364 19.885319
257 36.079083
491 6.343983
464 21.439839
11 20.642092
220 29.352433
336 21.041362
317 18.284488
218 27.475722
392 16.519279
199 34.940209
111 20.800415
32 18.696790
45 23.182934
143 16.590046
233 37.244496
297 27.297927
182 26.460052
408 18.002708
131 20.242091
204 39.798027
148 10.059654
288 26.590124
114 20.991302
211 22.979552
337 20.257692
37 21.764018
466 16.943828
443 17.902941
366 11.516250
151 21.521131
145 10.177929

```


## 206 22.246113
## 385 3.142551
## 372 22.572811
## 168 22.800943
## 333 24.122567
## 411 14.471592
## 481 21.790508
## 303 28.214483
## 170 26.147989
## 138 19.599097
## 320 21.606944
## 86 27.491503
## 299 29.190184
## 469 16.125460
## 485 18.346520
## 60 21.066863
## 100 32.060089
## 25 16.634002
## 266 26.754371
## 362 18.286173
## 321 24.705911
## 328 20.028764
## 187 35.029155
## 429 14.165937
## 331 23.325692
## 247 21.057087

#Color Vector
colors <- rainbow(n=length(Boston))

#plot lasso coefficients by lambda value
plot(x=as.vector(unlist(cvfit$glmnet.fit[5])), y=as.vector(coef(cvfit$glmnet.fit)[2,]),
      xlab = "Lambda", ylab="Coefficient Size", main="Coefficients Changing by Lambda - LASSO",
      type = "b",ylim=c(-18,5),col=colors[1])
for (i in 3:length(Boston)) {
  lines(x=as.vector(unlist(cvfit$glmnet.fit[5])), y=as.vector(coef(cvfit$glmnet.fit)[i,]),
        type = "b",ylim=c(-20,30),col=colors[i])
}

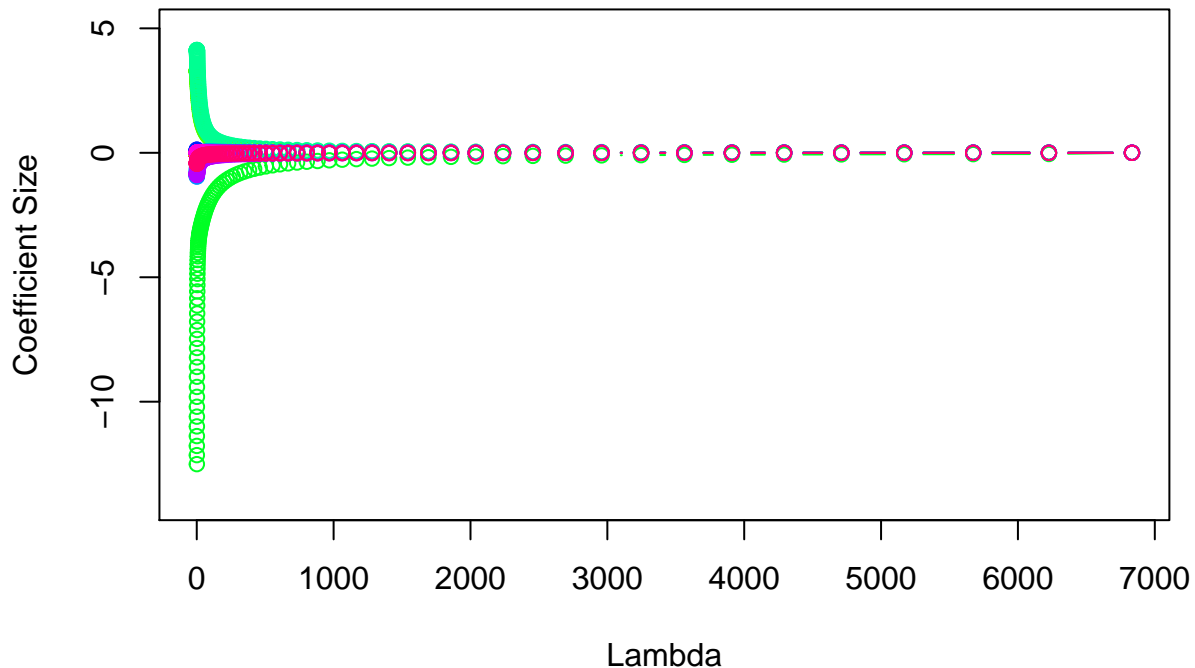
```

Coefficients Changing by Lambda – LASSO



```
#plot ridge coefficients by lambda value
plot(as.vector(unlist(cvfit2$glmnet.fit[5])), as.vector(coef(cvfit2$glmnet.fit)[2,]),
      xlab = "Lambda", ylab="Coefficient Size", main="Coefficients Changing by Lambda - Ridge",
      type = "b",ylim=c(-14,5),col=colors[1])
for (i in 3:length(Boston)) {
  lines(as.vector(unlist(cvfit2$glmnet.fit[5])), as.vector(coef(cvfit2$glmnet.fit)[i,]),
        type = "b",col=colors[i])
}
```

Coefficients Changing by Lambda – Ridge



Comments: As lambda increases, our coefficients all shrink in size. In ridge regression, the coefficients decrease in more curved shapes and asymptotically approach zero. With lasso, coefficients decrease in a spline shape and some coefficients are set to zero. This is because lasso does coefficient by using the ℓ_1 norm of the beta coefficients.

Problem 3 (25 points)

In this exercise, we will predict the number of applications received using the other variables in the `College` data set in the ISLR package.

Part a.

Randomly sample 20 percent of the rows of the dataset, and set this aside as a test set. Let the remainder be the training set.

```
#Lets convert the private variable to a numeric binary indicator.
College$Private <- ifelse(College$Private=="Yes",1,0)

set.seed(365)
testrows <- sample(nrow(College),round(nrow(College)*.20))
test <- College[testrows,]
train <- College[-testrows,]
```

Part b.

Fit a linear model using least squares on the training set, and report the test error obtained.

```

set.seed(365)
mod1 <- lm(Apps ~ ., data=train)
predapps <- predict(mod1, newdata=test[-2])
MSE_lm <- sum((test[2]-predapps)^2)/nrow(test)
MSE_lm
## [1] 1033568

```

Part c.

Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test error obtained.

```

set.seed(1)
cvmod2_ridge <- cv.glmnet(x=as.matrix(train[-2]), y=as.matrix(train[2]),
                          standardize=TRUE, nfold=10, alpha=0)
predapps_ridge <- predict(cvmod2_ridge, newx=as.matrix(test[-2]), s="lambda.min")
MSE_ridge <- sum((test[2]-predapps_ridge)^2)/nrow(test)
MSE_ridge
## [1] 885803.7

```

Part d.

Fit a lasso model on the training set, with λ chosen by cross-validation. Report the test error obtained, along with the number of non-zero coefficient estimates.

```

set.seed(1)
cvmod3_lasso <- cv.glmnet(x=as.matrix(train[-2]), y=as.matrix(train[2]),
                          standardize=TRUE, nfold=10, alpha=1)
predapps_lasso <- predict(cvmod3_lasso, newx=as.matrix(test[-2]), s="lambda.min")
MSE_lasso <- sum((test[2]-predapps_lasso)^2)/nrow(test)
MSE_lasso
## [1] 998732.2

```

Part e.

Comment on the results obtained. How accurately can we predict the number of college applications received? Is there much difference among the test errors resulting from these approaches? Which model do you think is better?

Comments: The mean squared error is highest for the linear model, decreases for the LASSO model and decreases more with the ridge regression model ($1,033,000 > 999,000 > 886,000$). No there is no a significant difference, based on a different seed we would get different values and thus likely a different order in which models produce the lowest MSE. According to `set.seed(1)` ridge regression is the best model.